

DS116 - Data Visualization

Intro to R

Habet Madoyan

American University of Armenia

Section 1

Intro to R

Intro to R

- Go to Comprehensive R Archive Network (CRAN) and install version for your operating system
- Go to RStudio and install RStudio IDE
- You need to install R before installing RStudio


Alternatives

- Install Microsoft R
- Use Microsoft Visual Studio (free version available)




Intro to R

Configure the pane layout of RStudio as you wish



Intro to R



Intro to R


RStudio projects make it straightforward to divide your work into multiple contexts, each with their own working directory, workspace, history, and source documents.

Creating Projects

- RStudio projects are associated with R working directories. You can create an RStudio project:
 - in a brand new directory;
 - in an existing directory where you already have R code and data;
 - by cloning a version control (Git or Subversion) repository.

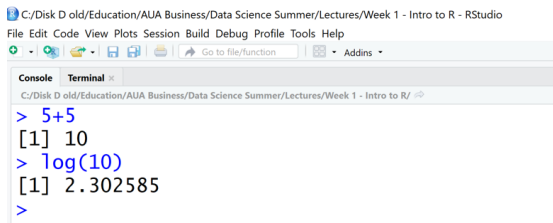
Intro to R

- Create an empty directory
- Go to File → New Project
- Choose existing directory



Intro to R

- Now when you have created your project, all files associated with the project need to be saved in the project directory
- To check the project directory do `getwd()` in Console



The screenshot shows the RStudio window with the project path `C:/Disk D old/Education/AUA Business/Data Science Summer/Lectures/Week 1 - Intro to R - RStudio` at the top. The menu bar includes File, Edit, Code, View, Plots, Session, Build, Debug, Profile, Tools, and Help. The toolbar contains icons for file operations and a search bar. The Console tab is active, displaying the following R commands and their outputs:

```
> 5+5
[1] 10
> log(10)
[1] 2.302585
>
```

The name of the project appears at the top of RStudio window

Why do we need RStudio Project?

The benefits of RStudio project:

- Can have different projects with their own environment.
- While opening a project, R restores previous work and project history includes recent commands which had been used in the project.
- Source pane remembers the files which had been opened.

Intro to R

When a new project is created, RStudio:

- Creates a project file (with an `.Rproj` extension) within the project directory. This file contains various project options (discussed below) and can also be used as a shortcut for opening the project directly from the filesystem.
- Creates a hidden directory (named `.Rproj.user`) where project-specific temporary files (e.g. auto-saved source documents, window-state, etc.) are stored. This directory is also automatically added to `.Rbuildignore`, `.gitignore`, etc. if required.
- Loads the project into RStudio and displays its name in the Projects toolbar (which is located on the far right side of the main toolbar).

Intro to R

Opening Projects


There are several ways to open a project:

- Using the **Open Project** command (available from both the Projects menu and the Projects toolbar) to browse for and select an existing project file (e.g. MyProject.Rproj).
- Selecting a project from the list of the most recently opened projects (also available from both the Projects menu and toolbar).
- Double-clicking on the project file within the system shell (e.g. Windows Explorer, OSX Finder, etc.).

Intro to R


The source editor

- If you plan to reuse your code, write it in source editor



Intro to R

Here you can type any valid R command after the `>` prompt followed by **Enter** and R will execute that command.

A screenshot of an R console window. The window has a title bar with 'Console' and 'Terminal' tabs. The path 'C:/Disk D old/Education/AUA Business/Data Science Summer/Lectures/Week 1 - Intro to R/' is visible. The prompt '>' is followed by the command '5+5'. The output is '[1] 10'. A new prompt '>' is shown on the next line with a cursor.


```
> 5+5
[1] 10
> |
```

Use console as a calculator

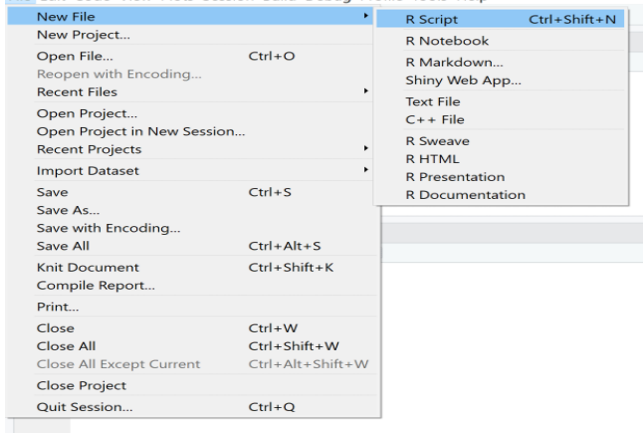
`[1]` is the index for the output, just ignore it

Intro to R

This will create new script file


 C:/Disk D old/Education/AUA Business/Data Science Summer/Lectures/Week 1 - Intro to R

File Edit Code View Plots Session Build Debug Profile Tools Help



Intro to R

- To run the script line from source editor, put the cursor anywhere on the line and hit **Ctrl+Enter**
- You will see the output in the console



```
class example.R
1
2 log(10)|
3
```

Intro to R

- If you have a piece of text in your editor that is not a code (thus is not executable by R) then you need to comment it, add `#` before each line
- When you run this line of code, it will be printed in Console as a text
- To comment large chunk of text, use **Ctrl+Shift+C**

C:\Disk D old\Education\AJA Business\Data Science Summer\Lectures\Week 1 - Intro to R/ >=


```
> # This command will calculate the logarithm of 10
>
> log(10)
[1] 2.302585
>
> # This line will calculate factorial of 10
>
> factorial(10)
[1] 3628800
>
> |
```

class example.R*

```
1
2 # This command will calculate the logarithm of 10
3
4 log(10)
5
6 # This line will calculate factorial of 10
7
8 factorial(10)
9
```


Intro to R

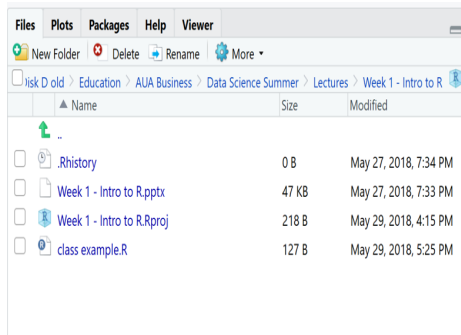
- The environment window contains objects (data, values, functions) R has currently stored in its memory.
- The history window shows all commands that were executed in the Console.



```
5+5  
log(10)  
log(10)  
log(10)  
log(10)  
factorial(10)  
factorial
```

Intro to R

Bottom right: files, plots, packages, help, & viewer pane. Here you can open files, view plots, install and load packages, read main pages, and view markdown and other documents in the viewer tab.



Intro to R

Some useful shortcuts:

- type first few letters of the function/object then hit Tab to open dropdown menu with possible options

C:/Disk D old/Education/AUA Business/Data Science Summer/Lectures/Week 1 - Intro to R - I

File Edit Code View Plots Session Build Debug Profile Tools Help

The screenshot shows the RStudio application window. The title bar indicates the file path: C:/Disk D old/Education/AUA Business/Data Science Summer/Lectures/Week 1 - Intro to R - I. The menu bar includes File, Edit, Code, View, Plots, Session, Build, Debug, Profile, Tools, and Help. Below the menu bar is a toolbar with icons for saving, opening, and navigating. The main window is divided into two panes: Console and Terminal. The Console pane is active and shows the command prompt with the text '> me'. A dropdown menu is displayed below the command prompt, listing five functions: mean, mean.Date, mean.default, mean.difftime, and mean.POSIXct. Each function is preceded by a small blue diamond icon and followed by '{base}' in a lighter font. The 'mean.POSIXct' option is currently selected and highlighted in blue. The Terminal pane is inactive and shows the same file path.


Function	Package
mean	{base}
mean.Date	{base}
mean.default	{base}
mean.difftime	{base}
mean.POSIXct	{base}

Intro to R

It's often the case that you want to re-execute commands that you previously entered. The RStudio console supports the ability to recall previous commands using the arrow keys:

- **Up** - Recall previous command(s)
- **Down** - Reverse of Up

You can even view a list of your recent commands by pressing **Ctrl+Up** on Windows or **Command+Up** on a Mac.

A screenshot of the RStudio console window. The title bar shows 'Console Terminal'. The address bar displays the file path 'C:/Disk D old/Education/AUA Business/Data Science Summer/Lectures/Week 1 - Intro to R/'. The console content shows a list of recent commands: 'factorial', 'gamma', 'source('C:/Disk D old/Education/AUA Business/Data Science Summer/Lecture...', '# This command will calculate the logarithm of 10', 'log(10)', '# This line will calculate factorial of 10', and 'factorial(10)'. A vertical scrollbar is visible on the right side of the console area.

```
factorial
gamma
source('C:/Disk D old/Education/AUA Business/Data Science Summer/Lecture:
# This command will calculate the logarithm of 10
log(10)
# This line will calculate factorial of 10
factorial(10)
```

Intro to R

- When you download R from the CRAN, you get that “base” R system.
- The base R system comes with basic functionality; implements the R language.
- One reason R is so useful is the large collection of packages that extend the basic functionality of R.
- R packages are developed and published by the larger R community.

Intro to R

- Packages can be installed with the `install.packages()` function in R.
- To install a single package, pass the name of the library to the `install.packages()` function as the first argument.
- The following code installs the `devtools` package from CRAN.

```
install.packages("devtools")
```

Intro to R

- The package needs to be installed only once
- To load the package into R environment you need to use function `library()`
- You need to load the library everytime you start a new R

```
library(devtools)
```

Intro to R

- You can also install packages from github
- The following code will install the following package from github

```
library(devtools)  
install_github("christophM/iml")
```

If you do not want to load the entire package but want to use some function from it, use the following command

package::function_name

```
devtools::install_github("christophM/iml")
```


Intro to R

To access help/documentation on a function from R base package

```
?mean
```

To access help/documentation on function from a library

```
###geom_path
```

The same

```
?ggplot2::geom_path
```

Help on the package

```
help(package='ggplot2')
```

Intro to R

- Each package on CRAN has its own webpage
- This includes documentation and sometimes includes vignettes
- If you have a specific task to do, then look at R Task View
- Here are all the packages and R functionality described for the Time series analysis

Section 2

Intro to R programming language

Intro to R programming language

- Anything in R is an object.
- Objects are assigned values using `<-` . (An equal sign `=` can also be used.)
For e.g., the following command assigns value 5 to object x.

```
x <- 5  
x  
## [1] 5
```

- R is case sensitive, thus Data10 and data10 are two different objects.
- A tidy code requires a space before the assignment operator and a space after.
 - You can see what is inside the object just by simple entering the name of the object in command line
 - When the object is created it should appear in your Environment Window
 - If the object is not in your environment window you cannot work with it
 - The assignment operator works in the opposite direction as well

```
5 -> x  
x  
## [1] 5
```

Data Structures in R


R programming supports five basic types of data structure.

- Homogeneous


Vector



Matrix




Array




- Heterogeneous

Data frame



List



Data Structures: Vector

Vector is a type of data structure that contains similar types of data, i.e., integer, double, logical, complex, etc. In order to create a vector `c()` function is used.

```
(x <- c(10,5,6))
```

```
## [1] 10 5 6
```

```
(x1 <- c(1:10))
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

The same as (only when you have a sequence)

```
(x1 <- 1:10)
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

```
class(x)
```

```
## [1] "numeric"
```

A character vector

```
y <- c("CS", "DS", "EC")
```

```
class(y)
```

```
## [1] "character"
```

Coercion

- All elements of an atomic vector must be the same type, so when you attempt to combine different types they will be coerced to the most flexible type.
- Types from least to most flexible are: logical, numeric, and character.
- For example, combining a character and an integer yields a character:

```
(x3 <- c("A",1))
```

```
## [1] "A" "1"
```

```
class(x3)
```

```
## [1] "character"
```

Data Structures: Matrix

- **Matrix** is a two-dimensional data structure and can be created using `matrix()` function.
- Matrix is a collection of vectors with the same length and type


```
m <- matrix(data=1:15, nrow=3)
```

```
m
```

```
##      [,1] [,2] [,3] [,4] [,5]  
## [1,]    1    4    7   10   13  
## [2,]    2    5    8   11   14  
## [3,]    3    6    9   12   15
```


Function

- To look at the function arguments, hit Tab
- You need to name the arguments if they are not in the same order as defined within the function
- If the order is the same you can skip the names



The screenshot shows an R console window with the following elements:

- Console** tab selected.
- Terminal path: `C:/Disk D old/Education/AUA Business/Data Science Summer/Lectures/Week 1 - Intro to R/`
- Command prompt: `> matrix()`
- A dropdown menu is open, listing the arguments: `data =`, `nrow =`, `ncol =`, `byrow =`, and `dimnames =`.
- A yellow tooltip box is visible on the right, titled **data**, containing the text: "an optional data vector (including a list or expression vector). Non-atomic classed R objects are coerced by `as.vector` and all attributes discarded." Below the text is a link: "Press F1 for additional help".

Data Structures

Example: function matrix with arguments flipped

The same result

```
(m <- matrix(1:15, 3))
```

```
##      [,1] [,2] [,3] [,4] [,5]  
## [1,]    1    4    7   10   13  
## [2,]    2    5    8   11   14  
## [3,]    3    6    9   12   15
```

Different result

```
(m <- matrix(3, 1:15))
```

```
##      [,1]  
## [1,]    3
```

Data Structures

Other ways of building a matrix

- Create two vectors of the same length

```
x <- 0:5  
y <- -5:0
```

- Combine together by row

```
(m1 <- rbind(x,y))
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6]  
## x      0   1   2   3   4   5  
## y     -5  -4  -3  -2  -1   0
```

Data Structures

The matrix has row names but does not have column names

```
colnames(m1)
```

```
## NULL
```

```
rownames(m1)
```

```
## [1] "x" "y"
```

Data Structures

- Combine together by column

```
(m2 <- cbind(x,y))
```

```
##      x  y  
## [1,] 0 -5  
## [2,] 1 -4  
## [3,] 2 -3  
## [4,] 3 -2  
## [5,] 4 -1  
## [6,] 5  0
```

Pay attention, the matrix has column names, but does not have row names

```
colnames(m2)
```

```
## [1] "x" "y"
```

```
rownames(m2)
```

```
## NULL
```

Data Structures

Check if the resulting object is a matrix

```
is.matrix(m1)
```

```
## [1] TRUE
```

```
is.matrix(m2)
```

```
## [1] TRUE
```

Data Structures

Set the rownames

```
rownames(m2) <- c('A', 'B', 'C', 'D', 'E', 'F')
```

```
m2
```

```
##   x  y  
## A 0 -5  
## B 1 -4  
## C 2 -3  
## D 3 -2  
## E 4 -1  
## F 5  0
```

Data Structures

Subset by either integer position, or by rowname

```
m2[1:3,]
```

```
##   x  y  
## A 0 -5  
## B 1 -4  
## C 2 -3
```

```
m2[c('A', 'B', 'C'),]
```

```
##   x  y  
## A 0 -5  
## B 1 -4  
## C 2 -3
```


Data Structures: Data frame

- A **data frame** is the most common way of storing data in R, and if used systematically makes data analysis easier.
- Under the hood, a data frame is a list of equal-length vectors that can also have different types.

```
df <- data.frame(x=1:10, y=11:20)  
head(df, n=5)
```

```
##   x  y  
## 1 1 11  
## 2 2 12  
## 3 3 13  
## 4 4 14  
## 5 5 15
```

Data Structures

- Dataframe always has both column and row names
- If you don't specify row names it is just the row number

```
colnames(df)
```

```
## [1] "x" "y"
```

```
rownames(df)
```

```
## [1] "1" "2" "3" "4" "5" "6" "7" "8" "9" "10"
```

Data Structures

- You can combine vectors with different type into a dataframe
- First create a matrix with two vectors of different types

```
Club <- c("Juventus", "Napoli", "Roma", "Inter")
Points <- c(95,91,77,72)
```

```
s_matrix <- cbind(Club, Points)
str(s_matrix)
```

```
## chr [1:4, 1:2] "Juventus" "Napoli" "Roma" "Inter" "95" "91" "77" "72"
## - attr(*, "dimnames")=List of 2
## ..$ : NULL
## ..$ : chr [1:2] "Club" "Points"
```

```
class(s_matrix)
```

```
## [1] "matrix" "array"
```

```
typeof(s_matrix)
```

```
## [1] "character"
```

Data Structures

Create a dataframe

```
df_seriea <- data.frame(Club, Points)  
str(df_seriea)
```

```
## 'data.frame':    4 obs. of  2 variables:  
##  $ Club   : chr  "Juventus" "Napoli" "Roma" "Inter"  
##  $ Points: num  95 91 77 72
```

Two columns (vectors) in the data frame have different types

Data Structures: List

List

- In R lists act as containers.
- Unlike vectors, the contents of a list are not restricted to a single type and can encompass any mixture of data types.
- Lists are sometimes called generic vectors, because the elements of a list can be of any type of R object, even lists containing further lists.
- This property makes them fundamentally different from atomic vectors.
- A list is a special type of vector. Each element can be a different type.

Data Structures

List containing dataframe and a matrix

```
my_list <- list(df_seriea, m1)
str(my_list)
```

```
## List of 2
## $ : 'data.frame': 4 obs. of 2 variables:
## ..$ Club : chr [1:4] "Juventus" "Napoli" "Roma" "Inter"
## ..$ Points: num [1:4] 95 91 77 72
## $ : int [1:2, 1:6] 0 -5 1 -4 2 -3 3 -2 4 -1 ...
## ..- attr(*, "dimnames")=List of 2
## .. ..$ : chr [1:2] "x" "y"
## .. ..$ : NULL
```

Data Structures

Look what is inside

```
my_list
```

```
## [[1]]
##      Club Points
## 1 Juventus    95
## 2   Napoli    91
## 3    Roma     77
## 4   Inter     72
##
## [[2]]
##      [,1] [,2] [,3] [,4] [,5] [,6]
## x      0   1   2   3   4   5
## y     -5  -4  -3  -2  -1   0
```

Data Structures

The `unlist()` function will flatten the list into a vector

```
z <- unlist(my_list)
```

```
z
```

	Club1	Club2	Club3	Club4	Points1	Points2	Points3
## "Juventus"		"Napoli"	"Roma"	"Inter"	"95"	"91"	"77"
## Points4							
## "72"		"0"	"-5"	"1"	"-4"	"2"	"-3"
##							
## "3"		"-2"	"4"	"-1"	"5"	"0"	

Data Structures

Attributes

- All objects can have arbitrary additional attributes, used to store metadata about the object.
- Attributes can be thought of as a named list (with unique names).
- Attributes can be accessed individually with `attr()` or all at once (as a list) with `attributes()`.

```
attr(dfSeriea, "topic") <- "sports"  
dfSeriea
```

```
##      Club Points  
## 1 Juventus    95  
## 2 Napoli     91  
## 3 Roma       77  
## 4 Inter      72
```

```
attr(dfSeriea, "topic")
```

```
## [1] "sports"
```

Section 3

Data Types

Data Types

Data types used in R

- Logical
- Numeric
- Factor
- Character

Data Types

Logical data type is one of the frequently used data types usually used for comparing two values. Values that a logical data type takes is TRUE or FALSE.

```
a <- 10  
a > 5
```

```
## [1] TRUE
```

```
log1 <- c(5, 6, TRUE)  
typeof(log1)
```

```
## [1] "double"
```

```
log1
```

```
## [1] 5 6 1
```

Coercion of numeric and logical values will result as numeric.

Data Types

String literals or string values are stored as Character objects in R.

```
b <- c("Armenia", "Georgia", "Azerbaijan")
typeof(b)
## [1] "character"
```

```
(b1 <- c(b, TRUE))
## [1] "Armenia"      "Georgia"      "Azerbaijan"  "TRUE"
```

Data Types

- A **factor** is a vector that can contain only predefined values, and is used to store categorical data.
- Factors are built on top of integer vectors using two attributes: the class, factor, which makes them behave differently from regular integer vectors, and the levels, which defines the set of allowed values.

```
b <- as.factor(c("Armenia", "Georgia", "Azerbaijan"))
```

Factors in R are stored as a vector of integer values with a corresponding set of character values used when the factor is displayed.

```
typeof(b)
```

```
## [1] "integer"
```

```
as.numeric(b)
```

```
## [1] 1 3 2
```

```
b
```

```
## [1] Armenia   Georgia   Azerbaijan
## Levels: Armenia Azerbaijan Georgia
```

Data Types

You can relevel the factor variable by changing the reference value

```
b<- relevel(b, ref='Azerbaijan')  
levels(b)
```

```
## [1] "Azerbaijan" "Armenia"      "Georgia"
```

```
as.numeric(b)
```

```
## [1] 2 3 1
```

Section 4

Special Values in R

Special Values in R

Missing Values

The examples and reasons of having missing values:

- The information was not collected (e.g. some people decline to give their age or weight).
- Some attributes are not applicable to all objects (e.g. forms have conditional parts that are filled out only when a person answers a previous question in a certain way).
- The information is not imported in the database (system missing).
- In R missing or undefined values are represented by NA.

Special Values in R

NA

- In R, the NA values are used to represent missing values. (NA stands for “not available”).
- You may encounter NA values in text loaded into R (to represent missing values) or in data loaded from databases (to replace NULL values).

```
v1 <- c(1,2,4,NA,5)
```

```
is.na(v1)
```

```
## [1] FALSE FALSE FALSE TRUE FALSE
```

Special Values in R

```
v2 <- c(10, "A", 20)
```

```
as.numeric(v2)
```

```
## Warning: NAs introduced by coercion  
## [1] 10 NA 20
```

Special Values in R

Inf and -Inf

If a computation results in a number that is too big, R will return Inf for a positive number and -Inf for a negative number (meaning positive and negative infinity, respectively).

```
120/0
```

```
## [1] Inf
```

```
-120/0
```

```
## [1] -Inf
```

Too big to show

```
45^12500
```

```
## [1] Inf
```

Special Values in R

NaN

Sometimes, a computation will produce a result that makes little sense. In these cases, R will often return NaN (meaning “Not a Number”).

```
0/0
```

```
## [1] NaN
```

Special Values in R

NULL

- There is a null object in R, represented by the symbol NULL.
- NULL is often used as an argument in functions to mean that no value was assigned to the argument.
- Some functions may return NULL.
- NULL is not the same as NA, Inf, -Inf, or NaN.
- NULL represents the null object in R.
- NULL is used mainly to represent the lists with zero length, and is often returned by expressions and functions whose value is undefined.

Section 5

Importing Data to R

Importing Data to R

Flat files

- A flat file database is a database that stores data in a plain text file.
- Each line of the text file holds one record with fields separated by delimiters, such as commas or tabs.
- While it uses a simple structure, a flat file database cannot contain multiple tables like a relational database can.

Importing Data to R

Delimiter-Separated Values

- Formats that use delimiter-separated values (also DSV) store two-dimensional arrays of data by separating the values in each row with specific delimiter characters.
- Most database and spreadsheet programs are able to read or save data in a delimited format.
- Most widely used delimiters are:
 - comma (CSV – comma separated values)
 - tab (TSV – tab separated values)

Importing Data to R

- File winter.csv contains data on winter Olympic games from 1924 to 2014.
- The file is comma-separated.
- Use readLines() with n=5 to look at the first five lines of the text file.
- You can see that the values are separated by comma.

```
readLines('Data/winter.csv', n=5)
```

```
## [1] "Year,City,Sport,Discipline,Athlete,Country,Gender,Event,Medal"
## [2] "1924,Chamonix,Biathlon,Biathlon,\"BERTHET, G.\",FRA,Men,Military Patrol,Bronze"
## [3] "1924,Chamonix,Biathlon,Biathlon,\"MANDRILLON, C.\",FRA,Men,Military Patrol,Bronze"
## [4] "1924,Chamonix,Biathlon,Biathlon,\"MANDRILLON, Maurice\",FRA,Men,Military Patrol,Bronze"
## [5] "1924,Chamonix,Biathlon,Biathlon,\"VANDELLE, André\",FRA,Men,Military Patrol,Bronze"
```

Importing Data to R

Use the function `read.csv()` to load the file into R environment

```
winter <- read.csv('Data/winter.csv')
str(winter)
```

```
## 'data.frame':    5770 obs. of  9 variables:
## $ Year      : int  1924 1924 1924 1924 1924 1924 1924 1924 1924 1924 ...
## $ City       : chr  "Chamonix" "Chamonix" "Chamonix" "Chamonix" ...
## $ Sport      : chr  "Biathlon" "Biathlon" "Biathlon" "Biathlon" ...
## $ Discipline: chr  "Biathlon" "Biathlon" "Biathlon" "Biathlon" ...
## $ Athlete    : chr  "BERTHET, G." "MANDRILLON, C." "MANDRILLON, Maurice" "VANDELLE, André" ...
## $ Country    : chr  "FRA" "FRA" "FRA" "FRA" ...
## $ Gender     : chr  "Men" "Men" "Men" "Men" ...
## $ Event      : chr  "Military Patrol" "Military Patrol" "Military Patrol" "Military Patrol" ...
## $ Medal      : chr  "Bronze" "Bronze" "Bronze" "Bronze" ...
```

Be sure that you can see the new object in the environment. Otherwise you cannot work with it



Importing Data to R

- Look the help for `read.csv` (`?read.csv`), several options on how the file is imported.
- `stringsAsFactors = F` – the strings are loaded as a text rather than as a factor.

```
winter <- read.csv('Data/winter.csv', stringsAsFactors = F)
str(winter)
```

```
## 'data.frame':    5770 obs. of  9 variables:
## $ Year      : int  1924 1924 1924 1924 1924 1924 1924 1924 1924 1924 ...
## $ City      : chr  "Chamonix" "Chamonix" "Chamonix" "Chamonix" ...
## $ Sport     : chr  "Biathlon" "Biathlon" "Biathlon" "Biathlon" ...
## $ Discipline: chr  "Biathlon" "Biathlon" "Biathlon" "Biathlon" ...
## $ Athlete   : chr  "BERTHET, G." "MANDRILLON, C." "MANDRILLON, Maurice" "VANDELLE, André" ...
## $ Country   : chr  "FRA" "FRA" "FRA" "FRA" ...
## $ Gender    : chr  "Men" "Men" "Men" "Men" ...
## $ Event     : chr  "Military Patrol" "Military Patrol" "Military Patrol" "Military Patrol" ...
## $ Medal     : chr  "Bronze" "Bronze" "Bronze" "Bronze" ...
```

Importing Data to R

- The same data is saved in tab delimited file.
- A tab-separated values (TSV) file is a simple text format for storing data in a tabular structure, e.g., database table or spreadsheet data, and a way of exchanging information between databases. Each record in the table is one line of the text file.

```
readLines('Data/winter.txt', n=5)
```

```
## [1] "\"Year\"\\t\"City\"\\t\"Sport\"\\t\"Discipline\"\\t\"Athlete\"\\t\"Country\"\\t\"Gender\"\\t\"Event\"\\t\"Medal\""
```

```
## [2] "\"1\"\\t1924\\t\"Chamonix\"\\t\"Biathlon\"\\t\"Biathlon\"\\t\"BERTHET, G.\"\\t\"FRA\"\\t\"Men\"\\t\"Military Patrol\"\\t\"Bronze\""
```

```
## [3] "\"2\"\\t1924\\t\"Chamonix\"\\t\"Biathlon\"\\t\"Biathlon\"\\t\"MANDRILLON, C.\"\\t\"FRA\"\\t\"Men\"\\t\"Military Patrol\"\\t\"Bronze\""
```

```
## [4] "\"3\"\\t1924\\t\"Chamonix\"\\t\"Biathlon\"\\t\"Biathlon\"\\t\"MANDRILLON, Maurice\"\\t\"FRA\"\\t\"Men\"\\t\"Military Patrol\"\\t\"Bronze\""
```

```
## [5] "\"4\"\\t1924\\t\"Chamonix\"\\t\"Biathlon\"\\t\"Biathlon\"\\t\"VANDELLE, André\"\\t\"FRA\"\\t\"Men\"\\t\"Military Patrol\"\\t\"Bronze\""
```

Importing Data to R

You can use the same `read.csv()` file but need to specify the separator (delimiter)

```
winter <- read.csv('Data/winter.txt', sep="\t", stringsAsFactors = FALSE)
str(winter)
```

```
## 'data.frame':   5770 obs. of  9 variables:
##  $ Year      : int  1924 1924 1924 1924 1924 1924 1924 1924 1924 1924 ...
##  $ City       : chr  "Chamonix" "Chamonix" "Chamonix" "Chamonix" ...
##  $ Sport      : chr  "Biathlon" "Biathlon" "Biathlon" "Biathlon" ...
##  $ Discipline: chr  "Biathlon" "Biathlon" "Biathlon" "Biathlon" ...
##  $ Athlete    : chr  "BERTHET, G." "MANDRILLON, C." "MANDRILLON, Maurice" "VANDELLE, André" ...
##  $ Country    : chr  "FRA" "FRA" "FRA" "FRA" ...
##  $ Gender     : chr  "Men" "Men" "Men" "Men" ...
##  $ Event      : chr  "Military Patrol" "Military Patrol" "Military Patrol" "Military Patrol" ...
##  $ Medal      : chr  "Bronze" "Bronze" "Bronze" "Bronze" ...
```

Importing Data to R

R can work with other data files as well

- STATA files
- SPSS
- SAS
- MS Excel
- Databases
- XML
- HTML
- Etc.

Section 6

Subsetting Data

Subsetting Data: Vectors

Create named vector

```
Club <- c("Juventus", "Napoli", "Roma", "Inter")
Points <- c(95,91,77,72)
names(Points) <- Club
Points
```

	Juventus	Napoli	Roma	Inter
	95	91	77	72

Subsetting Data: Vectors

- In R (unlike Python) the indexing starts with 1.
- Subsetting by index location.
- The vector has one dimension, so to subset it you need to specify the location of 1 index only.

```
Points[2]
```

```
## Napoli  
##      91
```

Subset by name

```
Points["Juventus"]
```

```
## Juventus  
##        95
```

Subsetting Data: Vectors

- Subset several elements
- First two elements

```
Points[1:2]
```

```
## Juventus Napoli
##      95      91
```

- 1 and 3 elements

```
Points[c(1,3)]
```

```
## Juventus      Roma
##      95      77
```

- Subsetting by name

```
Points[c("Napoli", "Inter")]
```

```
## Napoli Inter
##      91      72
```

Subsetting Data: Vectors

Why R you getting an eRRoR ?

```
Club[1,2]
```

```
## Error in Club[1, 2]: incorrect number of dimensions
```

```
Club[2,4]
```

```
## Error in Club[2, 4]: incorrect number of dimensions
```

Subsetting Data: Data frames

Load nba dataset

```
load('Data/nba2009_2018.rda')
summary(nba2009_2018)
```

```
## SEASON_ID GAME_DATE home.TEAM_ABBREVIATION
## Length:12060 Min. :2009-10-27 Length:12060
## Class :character 1st Qu.:2012-03-05 Class :character
## Mode :character Median :2014-11-13 Mode :character
## Mean :2014-08-10
## 3rd Qu.:2017-01-24
## Max. :2019-04-10
## home.TEAM_NAME home.PTS away.TEAM_ABBREVIATION away.TEAM_NAME
## Length:12060 Min. : 0.0 Length:12060 Length:12060
## Class :character 1st Qu.: 95.0 Class :character Class :character
## Mode :character Median :103.0 Mode :character Mode :character
## Mean :103.6
## 3rd Qu.:112.0
## Max. :161.0
## away.PTS home.WL
## Min. : 0.0 Length:12060
## 1st Qu.: 92.0 Class :character
## Median :100.0 Mode :character
## Mean :100.9
## 3rd Qu.:109.0
## Max. :168.0
```

Subsetting Data: Data frames

Data frame has two dimensions: Rows (first dimension) and Columns (second dimension)

`dataframe[,]`



Placeholder for rows

Placeholder for columns

Subsetting Data: Data frames

Will return the value on first row and forth column

```
dataframe[1, 4]
```

Will return the value of the first 10 rows and columns 4,5,6

```
dataframe[1:10, 4:6]
```

Subsetting Data: Data frames

Will return the first 10 rows and all columns

```
dataframe[1:10, ]
```

If you want to select all elements for the given index, just leave the placeholder empty.

What will this command return?

```
dataframe[, c(2,4:6)]
```


Subsetting Data: Data frames

Indexing by column names

```
colnames(nba2009_2018)
```

```
## [1] "SEASON_ID"      "GAME_DATE"      "home.TEAM_ABBREVIATION"
## [4] "home.TEAM_NAME" "home.PTS"       "away.TEAM_ABBREVIATION"
## [7] "away.TEAM_NAME" "away.PTS"       "home.WL"
```

```
nba1 <- nba2009_2018[,c("home.PTS", "away.PTS")]
```

```
colnames(nba1)
```

```
## [1] "home.PTS" "away.PTS"
```

```
dim(nba1)
```

```
## [1] 12060      2
```

Subsetting Data: Data frames

- Negative indexing is used to exclude certain records from the dataframe.
- This does not work with column names indexing.

```
nba2 <- nba2009_2018[, -c(1,2,4:6)]
head(nba2)
```

	home.TEAM_ABBREVIATION	away.TEAM_NAME	away.PTS	home.WL
## 1	CLE	Boston Celtics	95	L
## 2	DAL	Washington Wizards	102	L
## 3	POR	Houston Rockets	87	W
## 4	LAL	LA Clippers	92	W
## 5	ATL	Indiana Pacers	109	W
## 6	ORL	Philadelphia 76ers	106	W

Subsetting Data: Data frames

Exercises

- Create new data frame from nba dataset
- Include first 100 rows and columns 2, 3, 5
- Exclude rows 250, 300 to 350 and exclude column 5

Subsetting Data: Data frames

You can access specific column in data frame by using dollar sign

```
mean(nba2009_2018$home.PTS)
```

```
## [1] 103.6102
```

```
table(nba2009_2018$home.WL)
```

```
##  
##      L      W  
## 4949 7110
```

Subsetting Data: Data frames

Subset one variable

- First approach gives you a dataframe with one columns
- Second gives you a vector

```
home_pts <- nba2009_2018['home.PTS']  
str(home_pts)
```

```
## 'data.frame':    12060 obs. of  1 variable:  
##  $ home.PTS: num  89 91 96 99 120 120 101 92 115 74 ...
```

```
home_pts <- nba2009_2018[, 'home.PTS']  
str(home_pts)
```

```
##  num [1:12060] 89 91 96 99 120 120 101 92 115 74 ...
```

Subsetting Data: Data frames

Conditional indexing

- Create new dataframe with games only from season 2009
- We need all the rows where the value for SEASON_ID is 2009
- Note that the type for column SEASON_ID is character

```
nba4 <- nba2009_2018[nba2009_2018$SEASON_ID=='2009',]
```

Check if everything is done right

```
table(nba4$SEASON_ID)
```

```
##  
## 2009  
## 1230
```

Subsetting Data: Data frames

- Take only seasons 2009 and 2010
- As the SEASON_ID is character we will do the following

```
nba5 <- nba2009_2018[nba2009_2018$SEASON_ID %in% c("2009", "2010"),]  
table(nba5$SEASON_ID)
```

```
##  
## 2009 2010  
## 1230 1230
```

Subsetting Data: Data frames

- Workaround: make SEASON_ID numeric vector

```
nba2009_2018$SEASON_ID <- as.numeric(nba2009_2018$SEASON_ID)
nba5 <- nba2009_2018[nba2009_2018$SEASON_ID < 2011,]
table(nba5$SEASON_ID)
```

```
##
## 2009 2010
## 1230 1230
```


Subsetting Data: Data frames

Logical operators in R

Operator	Description
<	less than
<=	less than or equal to
>	greater than
>=	greater than or equal to
==	exactly equal to
!=	not equal to
!x	Not x
x y	x OR y
x & y	x AND y

Subsetting Data: Data frames

- Indexing on multiple conditions
- Take all the home game records for Detroit Pistons for seasons 2010 and 2011

```
nba6 <- nba2009_2018[nba2009_2018$SEASON_ID %in% c(2010,2011) &
                    nba2009_2018$home.TEAM_NAME == "Detroit Pistons",]
```

```
table(nba6$home.TEAM_NAME, nba6$SEASON_ID)
```

```
##
##           2010 2011
## Detroit Pistons    41   33
```

Subsetting Data: Data frames

Using OR (|) operator in R

```
nba7 <- nba2009_2018[nba2009_2018$away.TEAM_NAME=="Detroit Pistons" |
                    nba2009_2018$home.TEAM_NAME == "Detroit Pistons",]
```

```
head(nba7)
```


##	SEASON_ID	GAME_DATE	home.TEAM_ABBREVIATION	home.TEAM_NAME	home.PTS
## 10	2009	2009-10-28	MEM	Memphis Grizzlies	74
## 23	2009	2009-10-30	DET	Detroit Pistons	83
## 36	2009	2009-10-31	MIL	Milwaukee Bucks	96
## 54	2009	2009-11-03	DET	Detroit Pistons	85
## 61	2009	2009-11-04	TOR	Toronto Raptors	110
## 74	2009	2009-11-06	ORL	Orlando Magic	110
##	away.TEAM_ABBREVIATION		away.TEAM_NAME	away.PTS	home.WL
## 10	DET		Detroit Pistons	96	L
## 23	OKC	Oklahoma City Thunder		91	L
## 36	DET		Detroit Pistons	85	W
## 54	ORL		Orlando Magic	80	W
## 61	DET		Detroit Pistons	99	W
## 74	DET		Detroit Pistons	103	W

Subsetting Data: Data frames

Adding new variable in dataframe

- Point differential

```
nba2009_2018$Diff <- nba2009_2018$home.PTS-nba2009_2018$away.PTS  
hist(nba2009_2018$Diff)
```



Indexing List

```
(seriea <- data.frame(Club, Points))
```

```
##           Club Points
## Juventus Juventus   95
## Napoli   Napoli   91
## Roma     Roma    77
## Inter    Inter    72
```

- We have created a dataframe. Note that the dataframe has rownames
- Create a list

```
(list1 <- list(Club, "Italy", 2017, seriea))
```

```
## [[1]]
## [1] "Juventus" "Napoli"   "Roma"     "Inter"
##
## [[2]]
## [1] "Italy"
##
## [[3]]
## [1] 2017
##
## [[4]]
##           Club Points
## Juventus Juventus   95
## Napoli   Napoli   91
## Roma     Roma    77
## Inter    Inter    72
```

Indexing List

- If you index the list with `[]`, the result is a list.
- If you index the list with `[[[]]]` the result has the same structure as the list element

```
list1[4]
```

```
## [[1]]
##           Club Points
## Juventus Juventus    95
## Napoli    Napoli    91
## Roma      Roma      77
## Inter     Inter     72
```

```
list1[[4]]
```

```
##           Club Points
## Juventus Juventus    95
## Napoli    Napoli    91
## Roma      Roma      77
## Inter     Inter     72
```

```
is.data.frame(list1[[4]])
```

```
## [1] TRUE
```

```
is.data.frame(list1[4])
```

```
## [1] FALSE
```

Indexing List

Two elements

```
list1[1:2]  
## [[1]]  
## [1] "Juventus" "Napoli"    "Roma"      "Inter"  
##  
## [[2]]  
## [1] "Italy"
```

Slicing element from the element

This will bring the second element from the first element in the list

```
list1[[1]][2]  
## [1] "Napoli"
```

Indexing List

Named elements in the list

```
list1 <- list(Teams = Club, country = "Italy",
              year = 2017, standings = seriea)
```

```
list1
```

```
## $Teams
## [1] "Juventus" "Napoli"   "Roma"     "Inter"
##
## $country
## [1] "Italy"
##
## $year
## [1] 2017
##
## $standings
##           Club Points
## Juventus Juventus   95
## Napoli   Napoli    91
## Roma     Roma      77
## Inter    Inter      72
```


Indexing List

Indexing by name

```
list1$standings
```

```
##           Club Points
## Juventus Juventus   95
## Napoli   Napoli    91
## Roma     Roma      77
## Inter    Inter     72
```

```
list1$year
```

```
## [1] 2017
```

Section 7

Functions

Functions

- In programming, you use functions to incorporate sets of instructions that you want to use repeatedly or that, because of their complexity, are better self-contained in a sub program and called when needed.
- A function is a piece of code written to carry out a specified task; it can or can not accept arguments or parameters and it can or can not return one or more values.
- In fact, there are several possible formal definitions of 'function' spanning from mathematics to computer science.
- **Generically, its arguments constitute the input and return values as output.**

Functions

- Functions have named arguments which potentially have default values.
- The formal arguments are the arguments included in the function definition.
- The formals function returns a list of all the formal arguments of a function.
- Not every function call in R makes use of all the formal arguments.
- Function arguments can be missing or might have default values.

Functions

The function `runif()` generates random numbers from continuous uniform distribution

Arguments are:

- `n`, number of observations, with no default - so needs to be specified
- `min`, the minimum
- `max`, the maximum

```
formals(runif)
```

```
## $n
##
##
## $min
## [1] 0
##
## $max
## [1] 1
```

```
args(runif)
```

```
## function (n, min = 0, max = 1)
## NULL
```

Functions

R functions' arguments can be matched **positionally** or **by name**. So the following calls to `runif()` are all equivalent (you can mix them as well).

- Matching by name

```
runif(min=1, n=10, max=2)
```

```
## [1] 1.120291 1.039831 1.040716 1.710022 1.647759 1.442891 1.491432 1.471562  
## [9] 1.741866 1.996648
```

- Matching by position

```
runif(10,1,2)
```

```
## [1] 1.267439 1.770249 1.640040 1.426882 1.050750 1.019916 1.773276 1.147958  
## [9] 1.733208 1.351730
```

Note, as the process is random you might get different results.

Functions

Matched by position, this will give an error

```
runif(1,10,2)
```

```
## Warning in runif(1, 10, 2): NAs produced  
## [1] NaN
```

Functions

- If the argument has a default value and is not defined then the default value is used.
- Generated 10 numbers with min=0, max=1

```
runif(10)
```

```
## [1] 0.882818504 0.169275823 0.657059403 0.137945563 0.616535203 0.008407116  
## [7] 0.115444378 0.709159585 0.393922719 0.049844107
```


Functions

You can look at the arguments of the function and their description by hitting **Tab** inside the brackets

```
> runif()
```



The screenshot shows the RStudio interface with the command `> runif()` entered in the console. A tooltip is displayed, showing the arguments for the `runif()` function. The arguments are listed in a table-like structure:

Argument	Description
<code>n</code>	number of observations. If <code>length(n) > 1</code> , the length is taken to be the number required.
<code>min</code>	
<code>max</code>	

Below the table, it says "Press F1 for additional help".

Functions

General rule: if you have to copy a code more than twice, write a function.

```
f <- function(a, b = 1, c = 2, d = NULL) {  
  what needs to be done  
}
```

In addition to not specifying a default value, you can also set an argument value to NULL.

Functions

- User-defined functions are stored in the global environment and can be accessed easily.
- The following function will calculate the x power of y (default value of y is 2).

```
foo <- function (x, y=2){  
  x^y  
}
```

```
foo(4)
```

```
## [1] 16
```

Functions

Arguments to functions are evaluated lazily, so they are evaluated only as needed.

```
foo <- function (x,y){  
  x^2  
}  
  
foo(3)  
## [1] 9
```

This function never actually uses the argument *y*, so calling *f*(2) will not produce an error because the 2 gets positionally matched to *x*.

Functions

- The return value of a function is the last expression in the function body to be evaluated.
- It can be also specified with the function `return()`

Functions

The following function will calculate the z-score of a vector

```
norm <- function(x) {  
  return((x-mean(x))/sd(x))  
}
```

```
norm(mtcars$mpg)
```

```
## [1]  0.15088482  0.15088482  0.44954345  0.21725341 -0.23073453 -0.33028740  
## [7] -0.96078893  0.71501778  0.44954345 -0.14777380 -0.38006384 -0.61235388  
## [13] -0.46302456 -0.81145962 -1.60788262 -1.60788262 -0.89442035  2.04238943  
## [19]  1.71054652  2.29127162  0.23384555 -0.76168319 -0.81145962 -1.12671039  
## [25] -0.14777380  1.19619000  0.98049211  1.71054652 -0.71190675 -0.06481307  
## [31] -0.84464392  0.21725341
```

Section 8

Dates and Times in R

Dates in R

- R provides several options for dealing with date and date/time data.
- The built in function `as.Date()` handles dates (without times)
- Library `chron` handles dates and times, but does not control for time zones
- The `POSIXct` and `POSIXlt` classes allow for dates and times with control for time zones.
- The general rule for date/time data in R is to use the simplest technique possible.
- Thus, for date only data, `as.Date()` will usually be the best choice. If you need to handle dates and times without time-zone information, the `chron` library is a good choice; the `POSIX` classes are especially useful when time-zone manipulation is important.

Dates in R

- Except for the **POSIXlt** class, dates are stored internally as the number of days or seconds from some reference date.
- The function `as.Date()` stores the number of days passed from 1970-01-01

```
x <- as.Date("2018-06-18")  
class(x)  
## [1] "Date"
```

```
as.numeric(x)  
## [1] 17700
```

```
as.numeric(as.Date("1970-01-01"))  
## [1] 0
```

Dates in R

- The date can come in different formats.
- Usually, you need to tell R in which format the date is.

Symbol	Meaning	Example
%d	day as a number (0-31)	01-31
%a	abbreviated weekday	Mon
%A	unabbreviated weekday	Monday
%m	month (00-12)	00-12
%b	abbreviated month	Jan
%B	unabbreviated month	January
%y	2-digit year	07
%Y	4-digit year	2007

```
as.Date("2018/01/15", format = "%Y/%m/%d")
```

```
## [1] "2018-01-15"
```

```
as.Date("01-15-2018", format = "%m-%d-%Y")
```

```
## [1] "2018-01-15"
```

Dates in R

```
oil <- read.csv("Data/oil.csv", stringsAsFactors = FALSE)
head(oil)
```

```
##      DATE    OPEN    HIGH    LOW  CLOSE    VOL
## 1 1-Oct-12 112.14 113.27 110.76 111.40 80055
## 2 2-Oct-12 111.40 111.70 110.55 110.55 29332
## 3 3-Oct-12 110.55 110.59 106.95 107.34 56307
## 4 4-Oct-12 107.44 111.79 107.24 111.36 61664
## 5 5-Oct-12 111.27 112.09 109.64 111.13 51704
## 6 7-Oct-12 111.11 111.13 111.02 111.02    57
```

```
str(oil)
```

```
## 'data.frame':    1643 obs. of  6 variables:
##  $ DATE : chr  "1-Oct-12" "2-Oct-12" "3-Oct-12" "4-Oct-12" ...
##  $ OPEN : num  112 111 111 107 111 ...
##  $ HIGH : num  113 112 111 112 112 ...
##  $ LOW  : num  111 111 107 107 110 ...
##  $ CLOSE: num  111 111 107 111 111 ...
##  $ VOL  : int  80055 29332 56307 61664 51704 57 40638 61077 68836 68708 ...
```

Dates in R

Format the DATE column with as.Date

```
oil$DATE<-as.Date(oil$DATE, format="%d-%b-%y")  
str(oil)
```

```
## 'data.frame':    1643 obs. of  6 variables:  
##  $ DATE : Date, format: "2012-10-01" "2012-10-02" ...  
##  $ OPEN : num  112 111 111 107 111 ...  
##  $ HIGH : num  113 112 111 112 112 ...  
##  $ LOW  : num  111 111 107 107 110 ...  
##  $ CLOSE: num  111 111 107 111 111 ...  
##  $ VOL  : int  80055 29332 56307 61664 51704 57 40638 61077 68836 68708 ...
```

Dates in R

- Now we can subset the data frame with DATE column.
- The following code is going to create a data frame with cases only from 2014

```
oil_2014 <- oil[oil$DATE >= "2014-01-01" & oil$DATE < "2015-01-01",]
head(oil_2014)
```

```
##          DATE  OPEN  HIGH   LOW  CLOSE  VOL
## 389 2014-01-02 110.74 111.07 107.49 107.65 32977
## 390 2014-01-03 107.66 108.42 106.44 106.65 35605
## 391 2014-01-06 106.68 107.60 106.28 106.67 35795
## 392 2014-01-07 106.73 107.33 106.56 107.01 30686
## 393 2014-01-08 107.08 107.54 106.60 106.83 36728
## 394 2014-01-09 106.83 107.71 105.70 106.00 55777
```

Dates in R

You can do arithmetic operations with class Date

```
as.Date("2018-06-18")+1
```

```
## [1] "2018-06-19"
```

```
as.Date("2018-06-18")-as.Date("2018-05-18")
```

```
## Time difference of 31 days
```

Dates in R

Extracting weekday

```
oil$WEEKDAY <- weekdays(oil$DATE)
head(oil)
```

##	DATE	OPEN	HIGH	LOW	CLOSE	VOL	WEEKDAY
## 1	2012-10-01	112.14	113.27	110.76	111.40	80055	Monday
## 2	2012-10-02	111.40	111.70	110.55	110.55	29332	Tuesday
## 3	2012-10-03	110.55	110.59	106.95	107.34	56307	Wednesday
## 4	2012-10-04	107.44	111.79	107.24	111.36	61664	Thursday
## 5	2012-10-05	111.27	112.09	109.64	111.13	51704	Friday
## 6	2012-10-07	111.11	111.13	111.02	111.02	57	Sunday

Dates in R

Extract month

```
oil$MONTH <- months(oil$DATE)
head(oil)
```

##	DATE	OPEN	HIGH	LOW	CLOSE	VOL	WEEKDAY	MONTH
## 1	2012-10-01	112.14	113.27	110.76	111.40	80055	Monday	October
## 2	2012-10-02	111.40	111.70	110.55	110.55	29332	Tuesday	October
## 3	2012-10-03	110.55	110.59	106.95	107.34	56307	Wednesday	October
## 4	2012-10-04	107.44	111.79	107.24	111.36	61664	Thursday	October
## 5	2012-10-05	111.27	112.09	109.64	111.13	51704	Friday	October
## 6	2012-10-07	111.11	111.13	111.02	111.02	57	Sunday	October

Dates in R

- To extract the day of the month, use `format()`
- The returned column is character, but you can make it numeric

```
oil$DAY <- format(oil$DATE, "%d")
str(oil)
```

```
## 'data.frame':   1643 obs. of  9 variables:
## $ DATE      : Date, format: "2012-10-01" "2012-10-02" ...
## $ OPEN      : num  112 111 111 107 111 ...
## $ HIGH      : num  113 112 111 112 112 ...
## $ LOW       : num  111 111 107 107 110 ...
## $ CLOSE     : num  111 111 107 111 111 ...
## $ VOL       : int  80055 29332 56307 61664 51704 57 40638 61077 68836 68708 ...
## $ WEEKDAY   : chr   "Monday" "Tuesday" "Wednesday" "Thursday" ...
## $ MONTH     : chr   "October" "October" "October" "October" ...
## $ DAY       : chr   "01" "02" "03" "04" ...
```

Dates in R

Make the variable numeric

```
oil$DAY <- as.numeric(oil$DAY)
head(oil)
```

##	DATE	OPEN	HIGH	LOW	CLOSE	VOL	WEEKDAY	MONTH	DAY
## 1	2012-10-01	112.14	113.27	110.76	111.40	80055	Monday	October	1
## 2	2012-10-02	111.40	111.70	110.55	110.55	29332	Tuesday	October	2
## 3	2012-10-03	110.55	110.59	106.95	107.34	56307	Wednesday	October	3
## 4	2012-10-04	107.44	111.79	107.24	111.36	61664	Thursday	October	4
## 5	2012-10-05	111.27	112.09	109.64	111.13	51704	Friday	October	5
## 6	2012-10-07	111.11	111.13	111.02	111.02	57	Sunday	October	7

Section 9

R Markdown


R Markdown

- R Markdown documents are fully reproducible.
- It allows to combine text and code together in one report.
- You can use multiple languages within markdown file (R, python, SQL).
- R Markdown supports dozens of static and dynamic output formats including HTML, PDF, MS Word, Beamer, HTML5 slides, etc.
- R markdown is run in its own environment.




R Markdown

The Workflow



R Markdown

To open new markdown file



R Markdown

The document below is a template R Markdown document. It includes the most familiar parts of an R Markdown document:

- A YAML header that contains some metadata (1)
- Narrative text written in Markdown (2)
- R code chunks surrounded by triple backticks `{r}` and followed by triple backticks; a syntax that comes from the knitr package (3)

```

1 ---
2 title: "Example"
3 author: "Habet Madoyan"
4 date: "June 16, 2018"
5 output: pdf_document
6 ---
7
8 An R markdown document with some text and code for Data Science class
9
10
11 ```{r}
12 head(mtcars, n=5)
13 ```

```

R Markdown

Labeling and reusing code chunks

- Apart from the popular code chunk options you have learned by now, you can define even more things in the curly braces that follow the triple backticks.
- An interesting feature available in knitr is the labeling of code snippets. The code chunk below would be assigned the label `simple_sum`:

```
```{r simple_sum, results = 'hide'}  
2 + 2
```
```

- However, because the results option is equal to `hide`, no output is shown.

R Markdown

- You can embed R code into the text of your document with the `r` syntax. Be sure to include the lower case `r` in order for this to work properly.
- R Markdown will run the code and replace it with its result, which should be a piece of text, such as a character string or a number.
- For example, the line below uses embedded R code to create a complete sentence:

The factorial of four is ``r factorial(4)``.

- When you render the document the result will appear as:

The factorial of four is 24.

- Inline code provides a useful way to make your reports completely automated.

R Markdown

LaTeX equations

- You can also use the Markdown syntax to embed latex math equations into your reports. To embed an equation in its own centered equation block, surround the equation with two pairs of dollar signs like this,

$$1 + 1 = 2$$

- To embed an equation inline, surround it with a single pair of dollar signs, like this:

$1 + 1 = 2$

- You can use all of the standard latex math symbols to create attractive equations.

R Markdown

LaTeX formula example

$$\frac{n!}{k!(n-k)!} = \binom{n}{k}$$

In Tex

`\frac{n!}{k!(n-k)!} = \binom{n}{k}`

Tex script needs to be written in the text portion of the document

R Markdown

Lists in R Markdown

To make a bullet list in Markdown, place each item on a new line after an asterisk or hyphen and a space, like this:

```
* item 1  
* item 2  
* item 3
```

You can make an ordered list by placing each item on a new line after a number followed by a period followed by a space, like this

```
1. item 1  
2. item 2  
3. item 3
```

In each case, you need to place a blank line between the list and any paragraphs that come before it.

R Markdown

R code chunks

You can embed R code into your R Markdown report with the knitr syntax. To do this, surround your code with two lines: one that contains triple backticks `{r}` and the below one contains triple backticks. The result is a code chunk that looks like this:

```
```{r}  
some code
```
```

When you render the report, R will execute the code. If the code returns any results, R will add them to your report.

R Markdown

Customize R code chunks

- You can customize each R code chunk in your report by providing optional arguments after the `r` in “`{r}`”, which appears at the start of the code chunk. Let's look at one set of options.
- R functions sometimes return messages, warnings, and even error messages. By default, R Markdown will include these messages in your report. You can use the **message**, **warning** and **error** options to prevent R Markdown from displaying these. If any of the options are set to `FALSE`, R Markdown will not include the corresponding type of message in the output. Packages often generate messages when you first load them with `library()`.
- For example, R Markdown would ignore any messages or warnings generated by the chunk below.

```
```{r, warning = FALSE, message = FALSE}  
library(dplyr)
```
```

R Markdown

Three of the most popular chunk options are `echo`, `eval` and `results`.

- If **`echo = FALSE`**, R Markdown will not display the code in the final document (but it will still run the code and display its results unless told otherwise).
- If **`eval = FALSE`**, R Markdown will not run the code or include its results, (but it will still display the code unless told otherwise).
- If **`results = 'hide'`**, R Markdown will not display the results of the code (but it will still run the code and display the code itself unless told otherwise).

R Markdown

Other important stuff

Emphasis

```
*italic*  **bold**  
  
_italic_  __bold__
```

Headers

```
# Header 1  
  
## Header 2  
  
### Header 3
```

If you want to start from a new page in pdf file, put in text part of the markdown

`\newpage`

`\pagebreak`

Section 10

Data Transformation: dplyr

Grammar of data manipulation

dplyr is a grammar of data manipulation, providing a consistent set of verbs that help you solve the most common data manipulation challenges:

- **mutate()** adds new variables that are functions of existing variables
- **select()** picks variables based on their names.
- **filter()** picks cases based on their values.
- **summarise()** reduces multiple values down to a single summary.
- **arrange()** changes the ordering of the rows.

These all combine naturally with `group_by()` which allows you to perform any operation “by group”.

Grammar of data manipulation

- Not necessary, but dplyr works the best with pipe like operator from **magrittr** package
- `%>%` operator takes the object from its left-hand side and use it as an argument in the function on the right-hand side

`%>%`
magrittr

Ceci n'est pas un pipe.

Grammar of data manipulation

- The “pipe” operation is a handy tool to make your code more legible: `%>%`.

Key points:

- It takes the output of your previous operation and uses it as an input to your next operation.
- You can determine where the previous argument goes with the period symbol, `.`, which acts as a placeholder.
- Understand how to use it by replacing the pipe operation with “then” (in your mind, not in the code). For example, **`filter(data, ...)`** `%>%` **`select(...)`** filters first then selects columns from the output of filter.

Grammar of data manipulation

Filtering USA games only

```
summer <- read.csv("Data/summer.csv", stringsAsFactors = F)
```

```
summer_usa <- summer %>% filter(Country=="USA")
```

```
table(summer_usa$Country, summer_usa$Medal)
```

```
##
##      Bronze Gold Silver
## USA   1098 2235   1252
```

Grammar of data manipulation

Filter and group_by

```
summer %>%
  filter(Country %in% c("USA", "FRA", "GBR")) %>%
  group_by(Country) %>%
  summarise(Count=n())
```

```
## # A tibble: 3 x 2
##   Country Count
##   <chr>   <int>
## 1 FRA     1396
## 2 GBR     1720
## 3 USA     4585
```

- The result is a tibble
- Count=n() creates a new variable named Count with frequencies, n() calculates frequencies

Grammar of data manipulation

Number of medals by country (is this long or wide format?)

```
summer %>%
  filter(Country %in% c("USA", "FRA", "GBR")) %>%
  group_by(Country, Medal) %>%
  summarise(Count=n())
```

```
## # A tibble: 9 x 3
## # Groups:   Country [3]
##   Country Medal Count
##   <chr>    <chr> <int>
## 1 FRA      Bronze  497
## 2 FRA      Gold    408
## 3 FRA      Silver  491
## 4 GBR      Bronze  553
## 5 GBR      Gold    546
## 6 GBR      Silver  621
## 7 USA      Bronze 1098
## 8 USA      Gold   2235
## 9 USA      Silver 1252
```

Grammar of data manipulation

Using ggplot with dplyr

```
summer %>%
  filter(Country %in% c("USA", "FRA", "GBR")) %>%
  group_by(Country, Medal) %>%
  summarise(Count=n())
```

```
## # A tibble: 9 x 3
## # Groups:   Country [3]
##   Country Medal Count
##   <chr>    <chr> <int>
## 1 FRA      Bronze  497
## 2 FRA      Gold   408
## 3 FRA      Silver  491
## 4 GBR      Bronze  553
## 5 GBR      Gold   546
## 6 GBR      Silver  621
## 7 USA      Bronze 1098
## 8 USA      Gold   2235
## 9 USA      Silver 1252
```